

LEARNING TO COMMUNICATE: ENHANCING MULTI-AGENT COORDINATION WITH REINFORCEMENT LEARNING

Anonymous authors

Paper under double-blind review

// State the TL;DR of the paper, what we are trying to do, and its relevance.

ABSTRACT

This paper investigates integrating reinforcement learning (RL) into a multi-agent framework enhanced by a large language model to optimize communication and coordination among agents, which is crucial for complex tasks. Traditional methods often fall short in dynamic environments due to their inability to adapt communication protocols effectively, necessitating advancements in AI-driven systems. To tackle this, we design a reward-based system where agents receive feedback on communication effectiveness and task performance. Our key contributions include developing a simplified reward function that emphasizes task completion and communication efficiency and employing RL algorithms such as Q-learning and policy gradients to adapt and optimize agents' communication policies. We validate our approach through extensive experiments, showing significant improvements in system performance, task success rates, communication overhead, and learning efficiency over baseline models without RL-driven communication.

1 INTRODUCTION

The rapid advancement of artificial intelligence (AI) and machine learning has ushered in an era where autonomous agents are increasingly deployed in complex, dynamic environments. One significant challenge within multi-agent systems (MAS) is effective inter-agent communication and coordination, which are essential for achieving collective goals. Traditional methods often fall short in dynamic settings, highlighting the necessity for more adaptive solutions (?).

Optimizing communication among agents is inherently difficult due to the dynamic and unpredictable nature of the environments in which these systems operate. Effective coordination requires agents to not only share information but also interpret and respond to it in ways that align with overarching objectives. Inefficient communication protocols can severely undermine overall system performance.

To address these challenges, this paper proposes a novel approach by integrating reinforcement learning (RL) mechanisms into a multi-agent framework driven by a large language model. By leveraging RL, agents can dynamically optimize their communication strategies, leading to enhanced coordination and performance in varying contexts.

Our key contributions are as follows:

- Integration of RL mechanisms to enhance inter-agent communication and coordination.
- Development of a reward-based system for feedback based on communication effectiveness and task performance.
- Creation of a simplified reward function that emphasizes successful task completion and communication efficiency.
- Implementation of RL algorithms, such as Q-learning and policy gradients, to update agent communication policies.
- Comprehensive evaluation of system performance, task success rates, communication overhead, and learning efficiency.

Our approach involves designing a reward-based system where agents receive feedback on their communication and task performance. Utilizing RL algorithms, agents continually update their communication policies to adapt to changing environments, ensuring efficient and effective communication strategies.

We validate our approach through extensive experiments, comparing the performance of RL-driven communication models against baseline models that lack RL capabilities. Key metrics such as system performance, task success rates, communication overhead, and learning efficiency are analyzed to demonstrate the effectiveness of the proposed solution (?).

In summary, this work not only advances the state-of-the-art in multi-agent communication strategies but also lays the groundwork for future research. Future work could explore more complex reward functions, diverse environments, and the integration of additional learning paradigms to further enhance agent coordination and performance.